A Tutorial on Variable Neighborhood Search

Pierre Hansen

GERAD and HEC Montreal

Nenad Mladenović

GERAD and Mathematical Institute, SANU, Belgrade

July, 2003

Abstract

Variable Neighborhood Search (VNS) is a recent metaheuristic, or framework for building heuristics, which exploits systematically the idea of neighborhood change, both in the descent to local minima and in the escape from the valleys which contain them. In this tutorial we first present the ingredients of VNS, i.e., Variable Neighborhood Descent (VND) and Reduced VNS (RVNS) followed by the basic and then the general scheme of VNS itself which contain both of them. Extensions are presented, in particular Skewed VNS (SVNS) which enhances exploration of far away valleys and Variable Neighborhood Decomposition Search (VNDS), a two-level scheme for solution of large instances of various problems. In each case, we present the scheme, some illustrative examples and questions to be addressed in order to obtain an efficient implementation.

Keywords: metaheuristics, heuristics, descent, valley, Variable Neighborhood Search, decomposition.

Résumé

La Recherche à Voisinage Variable (RVV) est une métaheuristique récente, ou cadre général pour bâtir des heuristiques, qui exploite systématiquement l'idée de changement de voisinage tant dans la descente vers des optimas locaux que dans la sortie des vallées qui les contiennent. Dans cet article introductif nous présentons d'abord les ingrédients de la RVV, c'est-à-dire la Descente à Voisinage Variable (DVV) et la Recherche à Voisinage Variable Réduite (RVVR) suivis des schémas de base et général de la RVV. Des extensions sont ensuites présentées, en particulier la Recherche à Voisinage Variable Biaisée (RVVB) qui améliore l'exploration de vallées éloignées et la Recherche à Voisinage Variable avec Décomposition (RVVD), un schéma à deux niveaux pour la résolution d'instances de grande taille de divers problèmes. Dans chaque cas, on présente le schéma, un exemple et des questions à considérer pour une implémentation efficace.

Mots-clés: métaheuristique, heuristique, descente, vallée, Recherche à Voisinage Variable, décomposition.

1 Introduction

Let us consider a combinatorial or global optimization problem

$$\min f(x) \tag{1}$$

subject to

$$x \in X \tag{2}$$

where $f(x)$ is the *objective function* to be minimized and X the set of *feasible solutions*. A solution $x^* \in X$ is *optimal* if

$$f(x^*) \le f(x), \ \forall x \in X; \tag{3}$$

an *exact algorithm* for problem (1)-(2), if one exists, finds an optimal solution x^*, together with the proof of its optimality, or shows that there is no feasible solution, i.e., $X = \emptyset$. Moreover, in practice, the time to do so should be finite (and not too large); if one deals with a continuous function one must admit a degree of tolerance i.e., stop when a feasible solution x^* has been found such that

$$f(x^*) < f(x) + \varepsilon, \ \forall x \in X \tag{4}$$

or

$$\frac{f(x^*) - f(x)}{f(x^*)} < \varepsilon, \ \forall x \in X \tag{5}$$

for some small positive ε.

Numerous instances of problems of the form (1) (2), arising in Operational Research and other fields, are too large for an exact solution to be found in reasonable time. It is well-known from complexity theory (Garey and Johnson, 1979; Papadimitriou, 1994) that thousands of problems are *NP-hard*, that no algorithm with a number of steps polynomial in the size of the instances is known and that finding one for any such problem would entail obtaining one for any and all of them. Moreover, in some cases where a problem admits a polynomial algorithm, the power of this polynomial may be so large that realistic size instances cannot be solved in reasonable time in worst case, and sometimes also in average case or most of the time.

So one is often forced to resort to *heuristics*, which yield quickly an approximate solution, or sometimes an optimal solution but without proof of its optimality. Some of these heuristics have a worst-case guarantee, i.e., the solution x_h obtained satisfies

$$\frac{f(x_h) - f(x)}{f(x_h)} \le \varepsilon, \ \forall x \in X \tag{6}$$

for some ε, which is however rarely small. Moreover, this ε is usually much larger than the error observed in practice and may therefore be a bad guide in selecting a heuristic. In

addition to avoiding excessive computing time, heuristics address another problem: local optima. A local optimum x_L of (1)-(2) is such that

$$f(x_L) \leq f(x), \ \forall x \in N(x_L) \cap X \tag{7}$$

where $N(x_L)$ denotes a neighborhood of x_L (ways to define such a neighborhood will be discussed below). If there are many local minima, the range of values they span may be large. Moreover, the globally optimum value $f(x^*)$ may differ substantially from the average value of a local minimum, or even from the best such value among many, obtained by some simple heuristic (a phenomenon called by Baum (1986), the Tchebycheff catastrophe). There are, however, many ways to get out of local optima and, more precisely, the valleys which contain them (or set of solutions from which the descent method under consideration leads to them).

Metaheuristics are general framework to build heuristics for combinatorial and global optimization problems. They have been the subject of intensive research since Kirkpatrick, Gellatt and Vecchi (1983) proposed *Simulated Annealing* as a general scheme for building heuristics which get out of local minima. Several other metaheuristics were soon proposed. For discussion of the best-known of them the reader is referred to the books of surveys edited by Reeves (1993), Glover and Kochenberger (2003) as well as to the tutorials of the present volume. Some of the many successful applications of metaheuristics are also mentioned there.

Variable Neighborhood Search (VNS) (Mladenović and Hansen, 1997, Hansen and Mladenović 1999, 2001c, 2003) is a recent metaheuristic which exploits systematically the idea of neighborhood change, both in descent to local minima and in escape from the valleys which contain them. VNS exploits systematically the following observations:

Fact 1 *A local minimum with respect to one neighborhood structure is not necessary so for another;*

Fact 2 *A global minimum is a local minimum with respect to all possible neighborhood structures.*

Fact 3 *For many problems local minima with respect to one or several neighborhoods are relatively close to each other.*

This last observation, which is empirical, implies that a local optimum often provides some information about the global one. This may for instance be several variables with the same value in both. However, it is usually not known which ones are such. An organized study of the neighborhood of this local optimum is therefore in order, until a better one is found.

Unlike many other metaheuristics, the basic schemes of VNS and its extensions are simple and require few, and sometimes no parameters. Therefore in addition to providing

very good solutions, often in simpler ways than other methods, VNS gives insight into the reasons for such a performance, which in turn can lead to more efficient and sophisticated implementations.

The tutorial is organized as follows. In the next section, we examine the preliminary problem of gathering information about the problem under study, and evaluating it. In Section 3 the first ingredient of VNS, i.e., *Variable Neighborhood Descent* (VND), which is mostly or entirely deterministic, is studied. Section 4 is devoted to the second ingredient, i.e., *Reduced Variable Neighborhood Search* (RVNS), which is stochastic. Both ingredients are merged in the basic and the general VNS schemes, described in Section 5. Extensions are then considered. *Skewed Variable Neighborhood Search* (SVNS), which addresses the problem of getting out of very large valleys is discussed in Section 6. Very large instances of many problems cannot be solved globally in reasonable time; *Variable Neighborhood Decomposition Search* (VNDS) studied in Section 7 is a two-level scheme which merges VNS with successive approximation (including a two-level VNS). Various tools for analyzing in detail the performance of a VNS heuristic, and then streamlining it are presented in Section 8. They include *distance-to-target diagrams* and *valley profiles*. In each of these sections basic schemes, or tools, are illustrated by examples from papers by a variety of authors. Questions to be considered in order to get an efficient implementation of VNS are also systematically listed. Promising areas of research are outlined in Section 9. Brief conclusions complete the tutorial in Section 10. Sources of further informations are listed in an Appendix.

2   Preliminaries: Documentation

Once a problem of the form (1) (2) has been selected for study and approximate solution by VNS, a preliminary step is to gather in a thorough way the papers written about it or closely related problems. Note that this may be a difficult task as papers are often numerous, dispersed among many journals and volumes of proceedings and the problem may appear (usually under different names) in several fields. Tools such as the *ISI Web of Knowledge*, *NEC Research's Citeseer* or even general web browsers such as *Google* may prove to be very useful.

There are several reasons for studying the literature on the selected problem:

(i) *Evaluating its difficulty*: is it *NP-hard*? Is it *strongly NP-hard*? (and hence admits no fully polynomial approximation scheme). If it is in *P*, what is the complexity of the best-known exact algorithm, and is it sufficiently low for realistic instances to be solvable in reasonable time?

(ii) *Evaluating the performance of previous algorithms*: are there some instances of (preferably real-word) data for the problem available (e.g. at http://www.informs.org/Resources/Resources/Problem_Instances/)? What are the largest instances solved exactly?

(iii) *Evaluating the performance of previous heuristics*: which metaheuristics have been applied to this problem? What are the performances of the resulting heuristics, in terms of size of problems solved, error and computing time (assuming comparison among computing environments, if needed, can be done in a fairly realistic way)?

(iv) *What steps are used* in the heuristics already proposed? What are the corresponding neighborhoods of the current solution? Are codes for these heuristics available? Are codes for simple descent methods available?

Question (i)'s role is to help to assess the need for a VNS (or others) heuristic for the problem considered. Question (ii) and (iii) aim at obtaining a benchmark to evaluate the performance of the VNS heuristic when it will be designed and implemented: a good heuristic should obtain optimal solutions for most and preferably all instances solved by an exact algorithm (which suffers from the additional burden of having to prove optimality). Moreover, the new heuristic should do as well as previous ones on most or all instances and substantially better than them on quite a few instances to be viewed as a real progress (doing slightly better on a few instances is just not sufficient).

Question (iv) aims at providing ingredients for the VNS heuristic, notably in its VND component; it also inquires indirectly about directions not yet explored. Accessorily, it raises the question of possible re-use of software, which is reasonable for standard steps, e.g., a descent with Newton's method or a variant thereof.

3   Variable Neighborhood Descent

A *steepest descent* heuristic (known also as *best improvement* local search) consists of choosing an initial solution x, finding a direction of steepest descent from x, within a neighborhood $N(x)$, and moving to the minimum of $f(x)$ within $N(x)$ along that direction; if there is no direction of descent, the heuristic stops, and otherwise it is iterated. This set of rules is summarized in Figure 1.

Initialization.
Choose f, X, neighborhood structure $N(x)$, initial solution x;

Current step (Repeat).
      (1) Find $x' = argmin_{x \in N(x)} f(x)$;
      (2) If $f(x') < f(x)$ set $x' \leftarrow x''$ and iterate; otherwise, stop.

Figure 1. Steepest descent heuristic.

Observe that a neighborhood structure $N(x)$ is defined for all $x \in X$; in discrete optimization problems it usually consists of all vectors obtained from x by some simple modification, e.g. complementing one or two components of a 0-1 vector. Then, at each step, the neighborhood $N(x)$ of x is explored completely. As this may be time-consuming, an alternative

is to use the *first descent* heuristic. Vectors $x' \in N(x)$ are then enumerated systematically and a move is made as soon as a descent direction is found. This is summarized in Figure 2.

Initialization.
Choose f, X, neighborhood structure $N(x)$, initial solution x;

Current step (Repeat).
      (1) Find first solution $x' \in N(x)$;
      (2) If $f(x') > f(x)$, find next solution $x'' \in N(x)$; set $x' \leftarrow x''$ and iterate
      (2); otherwise, set $x \leftarrow x'$ and iterate (1);
      (3) If all solutions of $N(x)$ have been considered, stop.

Figure 2. First descent heuristic.

VND is based on *Fact 1* of the Introduction, i.e., *a local optimum for a first type of move* $x \leftarrow x'$ (or heuristic, or within the neighborhood $N_1(x)$) *is not necessary one for another type of move* $x \leftarrow \tilde{x}$ (within neighborhood $N_2(x)$)). It may thus be advantageous to combine descent heuristics. This leads to the basic VND scheme presented in Figure 3.

Initialization. Select the set of neighborhood structures N_ℓ, for $\ell = 1, \ldots, \ell_{max}$, that will be used in the descent; find an initial solution x (or apply the rules to a given x);

Repeat the following sequence until no improvement is obtained:
      (1) Set $\ell \leftarrow 1$;
      (2) *Repeat* the following steps until $\ell = \ell_{max}$:
            (a) Exploration of neighborhood.   Find the best neighbor x' of x ($x' \in$
            $N_\ell(x)$);
            (b) Move or not.   If the solution x' thus obtained is better than x, set
            $x \leftarrow x'$ and $\ell \leftarrow 1$; otherwise, set $\ell \leftarrow \ell + 1$;

Figure 3. Steps of the basic VND.

Caution should be exercised when applying that scheme. In particular one should consider the following questions:

(i) what complexity do the different moves have?

(ii) what is the best order in applying them?

(iii) are the moves considered sufficient to ensure a thorough exploration of the region containing x?

(iv) how precise a solution is desired?

Question (i) aims at selecting and ranking moves: if they involve too many elementary changes (e.g. complementing 3 components or more of a 0-1 vector), the resulting heuristic may be very slow and often take more time than an exact algorithm on small or medium size examples.

Question (ii) also bears upon computing times in relation to the quality of solutions obtained. A frequent implementation consists of ranking moves by order of complexity of their application (which is often synonymous with by size of their neighborhoods $|N_\ell(x)|$), and returning to the first one each time a direction of descent is found and a step made in that direction. Alternatively, all moves may be applied in sequence as long as descent is made for some neighborhood in the series.

Question (iii) is a crucial one: for some problems elementary moves are not sufficient to leave a narrow valley, and heuristics using them only can give very poor results. This is illustrated in Example 2 below.

Finally, the precision desired, as asked for in question (iv) will depend upon whether VND is used alone or within some larger framework, such as VNS itself. In the former case, one will strive to obtain the best solution possible within the allocated computing time; in the latter, one may prefer to get a good solution fairly quickly by the deterministic VND and to improve it later by faster stochastic search in VNS.

Example 1. Simple Plant Location (see e.g. Cornuejols, Fisher and Nemhauser, 1990, for a survey). The simple (or uncapacitated) plant location problem consists of locating a set of facilities i among a given set I of m potential locations, with fixed costs f_i, in order to minimize total costs for satisfying the demand of a given set of users J with delivery costs $c_{ij}, i \in I, j \in J$. It is expressed as follows:

$$\min_{x,y} z_P = \sum_{i=1}^{m} f_i y_i + \sum_{i=1}^{m} \sum_{j=1}^{n} c_{ij} x_{ij} \tag{8}$$

s.t.

$$\sum_{i=1}^{m} x_{ij} = 1, \ \forall j \in J; \tag{9}$$

$$y_i - x_{ij} \geq 0, \ \forall i \in I, \ \forall j \in J; \tag{10}$$

$$y_i \in \{0,1\}, \ \forall i \in I; \tag{11}$$

$$x_{ij} \geq 0, \ \forall i \in I, \ \forall j \in J. \tag{12}$$

where $y_i = 1$ if a facility is located at i, and 0 otherwise; $x_{ij} = 1$ if demand of user j is satisfied from facility i and 0 otherwise. Note that for fixed y_i, the best solution is defined by

$$x_{ij} = \begin{cases} 1 & \text{if } c_{ij} = \min_{\ell|y_\ell=1} c_{\ell j} \ \text{(with minimum index } \ell \text{ in case of ties)}; \\ 0 & \text{otherwise} \end{cases}$$

Therefore neighborhoods can be defined on the y_i, e.g. by Hamming distance (or number of components with complementary values). A first heuristic, *greedy*, proceeds by opening a facility ℓ with minimum total cost:

$$f_\ell + \sum_j c_{\ell j} = \min_i \{f_i + \sum_j c_{ij}\} \tag{13}$$

then letting

$$c_{rj} = \min_{i|y_i=1} c_{ij}, \ \forall j \tag{14}$$

computing the gains g_i obtained by opening a facility at i

$$g_i = \sum_j \max\{c_{rj} - c_{ij}, 0\} - f_i \tag{15}$$

and iteratively opening the facility for which the gain is larger, as long as it is positive. Each iteration takes $O(mn)$ time.

Once the greedy heuristic has been applied, an improved solution may be obtained by the *interchange heuristic* which proceeds iteratively to the relocation of one facility at a time in the most profitable way. With an efficient implementation, the idea of which was suggested by Whitaker (1983) for the closely related p-median problem, an iteration of interchange can also be made in $O(mn)$ time.

Applying in turn Greedy and Interchange is a simple case of VND. Further moves in which one facility would be closed and two opened, or two closed and one opened, or two opened and two closed would be too costly if all possible exchanges are examined.

Example 2. Minimum sum-of-squares clustering (MSSC): given N points $a_\ell \in \mathrm{I\!R}^p$ the minimum sum-of-squares clustering problem consists of partitioning them in M classes (or clusters) C_j such as to minimize the sum of squared distances between the points and the centroids $\overline{x}_i$ of their clusters:

$$\min \ \sum_{i=1}^m \sum_{\ell: a_\ell \in C_i} \|a_\ell - \overline{x}_i\|^2, \tag{16}$$

where

$$\overline{x}_i = \frac{1}{|C_i|} \sum_{\ell: a_\ell \in C_i} a_\ell, \tag{17}$$

and $\|.\|$ denotes the Euclidean norm.

Traditional heuristics for MSSC are: (i) H-Means, which proceeds from an initial partition by moving one entity x_ℓ from its cluster to another one, in a greedy way, until no further move decreases the objective function value, and (ii) K-Means, which proceeds from an initial partition by, alternatingly, finding the centroids of its clusters, and reassigning entities to the closest centroid, until stability is attained.

Computational experiments (Hansen and Mladenović, 2001) show that both H-Means and K-Means may lead to very poor results for instances with large M and N (the relative error being sometimes greater than 100%). This is due to bad exploration of X, or in other words, to difficulties in leaving valleys. A new "jump" move, defined as the displacement of a centroid to a point a_ℓ which does not coincide with a centroid, leads to a new VND heuristic, called J-Means, which improves very substantially on both H-Means and K-Means.

4 Reduced Variable Neighborhood Search

Assume a local minimum x of f has been reached. One would then like to leave its valley, and find another deeper one. In the standard versions of Variable Neighborhood Search, no previous knowledge of the landscape is assumed, or exploited. (Note that interesting hybrids could be built, using also values of $f(x)$ at previous iteration points x). Then, the questions to be asked are:

(i) in which direction to go?

(ii) how far?

(iii) how should one modify moves if they are not successful?

Question (i) bears upon the possibility of reaching any feasible point $x \in X$, or every valley; the simplest answer is to choose a direction at random. For problems in 0-1 variables this will amount to complementing some variables; for continuous Euclidean problems, drawing angular coefficients at random (or, in other words, choosing at random a point on the unit ball around x) takes all points of X into account.

Question (ii) is crucial. Indeed one wants to exploit to the limit *Fact 2* of the Introduction, i.e., in many combinatorial and global optimization problems, local optima tend to be close one to another and situated in one (or sometimes several) small parts of X. So once a local optimum has been reached, it contains implicit information about close better, and perhaps globally optimum, ones. It is then natural to explore first its vicinity. But, if the valley surrounding the local optimum x is large, this may not be sufficient, and what to do next is asked for in question (iii). Again a natural answer is to go further.

These aims are pursued in the *Reduced Variable Neighborhood* scheme, presented in Figure 4. A set of neighborhoods $N_1(x), N_2(x), \ldots, N_{kmax}(x)$ will be considered around the current point x (which may be or not a local optimum). Usually, these neighborhoods will be nested, i.e., each one contains the previous. Then a point is chosen at random in the first neighborhood. If its value is better than that of the incumbent (i.e., $f(x') < f(x)$), the search is recentered there ($x \leftarrow x'$). Otherwise, one proceeds to the next neighborhood. After all neighborhoods have been considered, one begins again with the first, until a stopping condition is satisfied (usually it will be maximum computing time since the last improvement, or maximum number of iterations).

Due to the nestedness property the size of successive neighborhoods will be increasing. Therefore one will explore more thoroughly close neighborhoods of x than farther ones, but nevertheless search within these when no further improvements are observed within the first, smaller ones.

Initialization.   Select the set of neighborhood structures $\mathcal{N}_k$, for $k = 1, \ldots, k_{max}$, that will be used in the search; find an initial solution x; choose a stopping condition;

Repeat the following sequence until the stopping condition is met:
  (1) Set $k \leftarrow 1$;
  (2) *Repeat* the following steps until $k = k_{max}$:
      (a) Shaking.   Generate a point x' at random from the k^{th} neighborhood of x $(x' \in \mathcal{N}_k(x))$;
      (b) Move or not.   If this point is better than the incumbent, move there $(x \leftarrow x')$, and continue the search with $\mathcal{N}_1$ $(k \leftarrow 1)$; otherwise, set $k \leftarrow k + 1$;

Figure 4. Steps of the Reduced VNS.

Example 3. *p*-**Median** (see e.g. Labbé, Peeters and Thisse, 1995 for a survey). This is a location problem very close to Simple Plant Location. The differences are that there are no fixed costs, and that the number of facilities to be opened is set at a given value p. It is expressed as follows:

$$\min \ \sum_{i=1}^{m}\sum_{j=1}^{n} c_{ij}x_{ij} \tag{18}$$

subject to

$$\sum_{i=1}^{m} x_{ij} \ = \ 1, \quad \forall j, \tag{19}$$

$$y_i - x_{ij} \ \geq \ 0, \quad \forall i,j, \tag{20}$$

$$\sum_{i=1}^{m} y_i \ = \ p, \tag{21}$$

$$x_{ij}, y_i \ \in \ \{0,1\}. \tag{22}$$

The Greedy and Interchange heuristics described above for Simple Plant Location are easily adapted to the p-Median problem and, in fact, the latter was early proposed by Teitz and Bart (1967).

Fast interchange, using Whitaker's (1983) data structure applies here also (Hansen and Mladenović, 1997). Refinements have recently been proposed by Resende and Werneck

(2002). A comparison between that approach and RVNS is made in Hansen *et al.* (2001), and results summarized in Table 1. It appears that RVNS gives better results than Fast Interchange in 40 times less time.

	OBJ. VALUE	CPU TIMES			% ERROR		
p	*Best known*	FI	RVNS	VNDS	FI	RVNS	VNDS
100	2733817.25	6637.48	510.20	6087.75	0.36	0.15	0.00
200	1809064.38	14966.05	663.69	14948.37	0.79	0.36	0.00
300	1394715.12	20127.91	541.76	17477.51	0.65	0.51	0.00
400	1145669.38	23630.95	618.62	22283.04	0.82	0.59	0.00
500	974275.31	29441.97	954.10	10979.77	0.98	0.51	0.00
700	752068.38	36159.45	768.84	32249.00	0.64	0.50	0.00
800	676846.12	38887.40	813.38	20371.81	0.61	0.53	0.00
900	613367.44	41607.78	731.71	27060.09	0.55	0.53	0.00
1000	558802.38	44176.27	742.70	26616.96	0.73	0.66	0.00
Average		28403.90	705.00	19786.00	0.68	0.48	0.00

Table 1: 5934-customer p-Median problem.

5 Basic and General Variable Neighborhood Search

In the previous two sections, we examined how to use variable neighborhoods in descent to a local optimum and in finding promising regions for near-optimal solutions. Merging the tools for both tasks leads to the General Variable Neighborhood Search scheme. We first discuss how to combine a local search with systematic changes of neighborhoods around the local optimum found. We then obtain the Basic VNS scheme presented in Figure 5.

Initialization.   Select the set of neighborhood structures $\mathcal{N}_k$, for $k = 1, \ldots, k_{max}$, that will be used in the search; find an initial solution x; choose a stopping condition;

Repeat the following sequence until the stopping condition is met:

    (1) Set $k \leftarrow 1$;

    (2) *Repeat* the following steps until $k = k_{max}$:

        (a) Shaking.   Generate a point x' at random from the k^{th} neighborhood of x ($x' \in \mathcal{N}_k(x)$);

        (b) Local search.   Apply some local search method with x' as initial solution; denote with x'' the so obtained local optimum;

        (c) Move or not.   If the local optimum x'' is better than the incumbent x, move there ($x \leftarrow x''$), and continue the search with $\mathcal{N}_1$ ($k \leftarrow 1$); otherwise, set $k \leftarrow k + 1$;

Figure 5. Steps of the basic VNS.

According to this basic scheme, a series of neighborhood structures, which define neighborhoods around any point $x \in X$ of the solution space, are first selected. Then the local search is used and leads to a local optimum x. A point x' is selected at random within the first neighborhood $\mathcal{N}_1(x)$ of x and a descent from x' is done with the local search routine. This leads to a new local minimum x''. At this point, three outcomes are possible: (i) $x'' = x$, i.e., one is again at the bottom of the same valley; in this case the procedure is iterated using the next neighborhood $\mathcal{N}_k(x)$, $k \geq 2$; (ii) $x'' \neq x$ but $f(x'') \geq f(x)$, i.e., another local optimum has been found, which is not better than the previous best solution (or incumbent); in this case too the procedure is iterated using the next neighborhood; (iii) $x'' \neq x$ and $f(x'') < f(x)$ i.e., another local optimum, better than the incumbent has been found; in this case the search is recentered around x'' and begins again with the first neighborhood. Should the last neighborhood be reached without a solution better than the incumbent being found, the search begins again at the first neighborhood $\mathcal{N}_1(x)$ until a stopping condition, e.g., a maximum time or maximum number of iterations or maximum number of iterations since the last improvement, is satisfied.

If instead of simple local search, one uses Variable Neighborhood Descent and if one improves the initial solution found by Reduced VNS, one obtains the General Variable Neighborhood Search scheme. This scheme is presented in Figure 6.

Initialization. Select the set of neighborhood structures $\mathcal{N}_k$, for $k = 1, \ldots, k_{max}$, that will be used in the shaking phase, and the set of neighborhood structures N_ℓ for $\ell = 1, \ldots, \ell_{max}$ that will be used in the local search; find an initial solution x and improve it by using RVNS; choose a stopping condition;

Repeat the following sequence until the stopping condition is met:
  (1) Set $k \leftarrow 1$;
  (2) *Repeat* the following steps until $k = k_{max}$:
    (a) Shaking. Generate a point x' at random from the k^{th} neighborhood $\mathcal{N}_k(x)$ of x;
    (b) Local search by VND.
      (b1) Set $\ell \leftarrow 1$;
      (b2) Repeat the following steps until $\ell = \ell_{max}$;
        · *Exploration of neighborhood.* Find the best neighbor x'' of x' in $N_\ell(x')$;
        · *Move or not.* If $f(x'') < f(x')$ set $x' \leftarrow x''$ and $\ell \leftarrow 1$; otherwise set $\ell \leftarrow \ell + 1$;
    (c) Move or not. If this local optimum is better than the incumbent, move there $(x \leftarrow x'')$, and continue the search with $\mathcal{N}_1$ $(k \leftarrow 1)$; otherwise, set $k \leftarrow k + 1$;

Figure 6. Steps of the General VNS.

Several questions about selection of neighborhood structures are in order:

(i) What properties of the neighborhoods are mandatory for the resulting scheme to be able to find a globally optimal or near-optimal solution?

(ii) What properties of the neighborhoods will be helpful in finding a near-optimal solution?

(iii) Should neighborhoods be nested? Otherwise how should they be ordered?

(iv) What are desirable properties of the sizes of neighborhoods?

The first two questions bear upon the ability of the VNS heuristic to find the best valleys, and to do so fairly quickly. To avoid being blocked in a valley, while there may be deeper ones, the union of the neighborhoods around any feasible solution x should contain the whole feasible set:

$$X \subseteq \mathcal{N}_1(x) \cup \mathcal{N}_2(x) \cup \cdots \cup \mathcal{N}_{k_{max}}(x), \ \ \forall x \in X.$$

These sets may cover X without necessarily partitioning it, which is easier to implement, e.g. when using nested neighborhoods, i.e.,

$$\mathcal{N}_1(x) \subset \mathcal{N}_2(x) \subset \cdots \subset \mathcal{N}_{k_{max}}(x), \ \ X \subset \mathcal{N}_{k_{max}}(x), \ \ \forall x \in X.$$

If these properties do not hold, one might still be able to explore X completely, by traversing small neighborhoods around points on some trajectory, but it is no more guaranteed. To illustrate, as mentioned before in minimum sum-of-squares clustering, the neighborhoods defined by moving an entity (or even a few entities) from one cluster to another one are insufficient to get out of many local optima. Moving centers of clusters does not pose a similar problem.

Nested neighborhoods are easily obtained for many combinatorial problems by defining a first neighborhood $\mathcal{N}_1(x)$ by a type of move (e.g. 2-opt in the traveling salesman problem) and then iterating it k times to obtain neighborhoods $\mathcal{N}_k(x)$ for $k = 2, \ldots, k_{max}$. They have the property that their sizes are increasing. Therefore if, as is often the case, one goes many times through the whole sequence of neighborhoods the first ones will be explored more thoroughly than the last ones. This is desirable in view of *Fact 3* mentioned in the Introduction, i.e., that local optima tend to be close one from another.

Restricting moves to the feasible set X may be too constraining, particularly if this set is disconnected. Introducing some or all constraints in the objective function with Lagrangian multipliers, allows moving to infeasible solutions. A variant of this idea is to penalize infeasibilities, e.g. pairs of adjacent vertices to which the same color is assigned in Graph coloring, see Zufferey, Hertz and Avanthay (2003).

Example 4. Scheduling workover rigs for onshore oil production.

Many oil wells in onshore fields rely on artificial lift methods. Maintenance services such as cleaning and others, which are essential to these wells, are performed by workover rigs. They are slow mobile units and, due to their high operation costs, there are relatively few workover rigs when compared with the number of wells demanding service. The problem of scheduling workover rigs consists in finding the best schedule S_i ($i = 1, \ldots, m$) of the m workover rigs to attend all wells demanding maintenance services, so as to minimize the oil production loss (production before maintenance being reduced).

In Aloise *et al.* (2003) a basic VNS heuristic is developed for solving the Scheduling of Workover Rigs Problem (WRP). Initial schedule S_i (where S_i is an ordered set of wells serviced by workover rig i), are obtained by a *Greedy* constructive heuristic. For the shaking step $k_{max} = 9$ neighborhoods are constructed: (1) *Swap routes* (SS): the wells and the associated routes assigned to two workover rigs are interchanged; (2) *Swap wells from the same workover rig* (SWSW): the order in which two wells are serviced by the same rig is swapped; (3) *Swap wells from different workover rig* (SWDW): two wells assigned to two different workover rigs are swapped; (4) *Add/Dropp* (AD): a well assigned to a workover rig is reassigned to any position of the schedule of another workover rig; (5) $(SWSW)^2$: apply twice the SWSW move; (6) $(SWDW)^2$: apply twice the SWDW move; (7) $(SWDW)^3$: apply three times the SWDW move; (8) $(AD)^2$: successively apply two (AD) moves; (9) $(AD)^3$: successively apply three (AD) moves.

For local search, the neighborhood consists of all possible exchanges of pairs of wells, i.e., the union of (SWSW) and (SWDW) from above are used.

A basic VNS is compared with Genetic algorithm, Greedy randomized Adaptive Procedure (GRASP) and two Ant colony methods (AS and MMAS) on synthetical and real life problems from Brasilian onshore fields. Some results on synthetical data are given in Table 2. On 27 possible scenarios in generating data sets (denoted by P-111, P-112, P-113, P-121,..., P333), VNS was better than others in 85% of the cases and MMAS in 15%. On real life problems, results were much better than the gains expected. For example, a daily reduction of 109 m^3 (equivalent to 685.6 bbl) in the production losses along 15 days was obtained by VNS compared with Petrobras' previous solution. That leads to a total savings estimated at 6,600,000 US dollars per year.

Problem	GA	GRASP	AS	MMAS	VNS
P-111	16791.87	16602.51	15813.53	15815.26	**15449.50**
P-211	20016.14	19726.06	19048.13	19051.61	**18580.64**
P-311	20251.93	20094.37	19528.93	19546.10	**19434.97**

Table 2: Average results with eight workover rigs over 20 runs of each synthetic test problem and three possible scenarios (from Aloise *et al.*, 2003).

6 Skewed Variable Neighborhood Search

VNS gives usually better or as good solutions than multistart, and much better ones when there are many local optima. This is due to *Fact 3* of the Introduction: many problems have clustered local optima; often, their objective function is a globally convex one plus some noise. However, it may happen that some instances have several separated and possibly far apart valleys containing near-optimal solutions. If one considers larger and larger neighborhoods, the information related to the currently best local optimum dissolves and VNS degenerates into multistart. Moreover if the current best local optimum is not

in the deepest valley this information is in part irrelevant. It is therefore of interest to modify VNS schemes in order to explore more fully valleys which are far away from the incumbent solution. This will be done by accepting to recenter the search when a solution close to the best one known, but not necessarily as good, is found, provided that it is far from this last solution. The modified VNS scheme for this variant, called *Skewed Variable Neighborhood Search* (SVNS) is presented in Figure 7. The relaxed rule for recentering uses an evaluation function linear in the distance from the incumbent, i.e., $f(x'')$ is replaced by

$$f(x'') - \alpha\rho(x, x'')$$

where $\rho(x, x'')$ is the distance from x to x'' and α a parameter. A metric for distance between solutions is usually easy to find, e.g. the Hamming distance when solutions are described by boolean vectors or the Euclidean distance in the continuous case.

Initialization.   Select the set of neighborhood structures $\mathcal{N}_k$, for $k = 1, \ldots, k_{max}$, that will be used in the search; find an initial solution x and its value $f(x)$; set $x_{opt} \leftarrow x$, $f_{opt} \leftarrow f(x)$: choose a stopping condition and a parameter value α;

Repeat the following until the stopping condition is met:
      (1) Set $k \leftarrow 1$;
      (2) *Repeat* the following steps until $k = k_{max}$:
        (a)*Shaking.*   Generate a point x' at random from the k^{th} neighborhood of x;
        (b)*Local search.*   Apply some local search method with x' as initial solution; denote with x'' the so obtained local optimum;
        (c) *Improvement or not.* If $f(x'') < f_{opt}$ set $f_{opt} \leftarrow f(x)$ and $x_{opt} \leftarrow x''$;
        (d) *Move or not.*   If $f(x'') - \alpha\rho(x, x'') < f(x)$ set $x \leftarrow x''$ and $k \leftarrow 1$; otherwise set $k \leftarrow k + 1$.

Figure 7. Steps of the Skewed VNS.

Clearly, more complicated formulas could be used for recentering; possibly, one might take into account known values at points already visited in the valley being explored.

Questions to be answered when applying SVNS are the following:

(i) does the problem under consideration have a roughly convex objective function, or are there several far apart deep valleys?

(ii) how should α be chosen?

These questions can be answered, to some extent, by first using a multistart version of VNS, i.e., starting VNS from various random points and running it for a short time. Then one can look at the position of the best local optima found and see if they are clustered or dispersed. Further, one can plot values in function of distance from the corresponding local optima to the best known solution and choose α as a fraction of the average slope.

Example 6. Weighted Maximum satisfiability (WMAX-SAT). The satisfiability problem, in clausal form, consists in determining if a given set of m clauses (all in disjunctive or all in conjunctive form) built upon n logical variables has a solution or not. The maximum satisfiability problem consists in finding a solution satisfying the largest possible number of clauses. In the *weighted maximum satisfiability* problem (WMAXSAT) positive weights are assigned to the clauses and a solution maximizing the sum of weights of satisfied clauses is sought. Results of comparative experiments with VNS and TS heuristics on instances having 500 variables, 4500 clauses and 3 variables per clause, in direct or complemented form, are given in Table 3 from Hansen *et al.* (2001). It appears that using a restricted neighborhood consisting of a few directions of steepest descent or mildest ascent in the Shaking step does not improve results, but using this idea in conjunction with SVNS improves notably upon results of basic VNS and also upon those of a TS heuristic.

	VNS	VNS-low	SVNS-low	TS
Number of instances where best solution is found	6	4	23	5
% Average error in 10 trials	0.2390	0.2702	0.0404	0.0630
% Best error in 10 trials	0.0969	0.1077	0.0001	0.0457
Total number of instances	25	25	25	25

Table 3: Results for *GERAD* test problems for WMAXSAT ($n = 500$).

7 Variable Neighborhood Decomposition Search

The **Variable Neighborhood Decomposition Search** (VNDS) method (Hansen *et al.*, 2001) extends the basic VNS into a two-level VNS scheme based upon decomposition of the problem. Its steps are presented on Figure 8.

Initialization. Select the set of neighborhood structures $\mathcal{N}_k$, for $k = 1, \ldots, k_{max}$, that will be used in the search; find an initial solution x; choose a stopping condition;

Repeat the following sequence until the stopping condition is met:
  (1) Set $k \leftarrow 1$;
  (2) *Repeat* the following steps until $k = k_{max}$:
    (a) Shaking. Generate a point x' at random from the k^{th} neighborhood of x ($x' \in \mathcal{N}_k(x)$); in other words, let y be a set of k solution attributes present in x' but not in x ($y = x' \setminus x$).
    (b) Local search. Find a local optimum in the space of y either by inspection or by some heuristic; denote the best solution found with y' and with x'' the corresponding solution in the whole space S ($x'' = (x' \setminus y) \cup y'$);
    (c) Move or not. If the solution thus obtained is better than the incumbent, move there ($x \leftarrow x''$), and continue the search with $\mathcal{N}_1$ ($k \leftarrow 1$); otherwise, set $k \leftarrow k + 1$;

Figure 8. Steps of the basic VNDS.

Note that the only difference between the basic VNS and VNDS is in step 2b: instead of applying some local search method in the whole solution space $\mathcal{S}$ (starting from $x' \in \mathcal{N}_k(x)$), in VNDS we solve at each iteration a subproblem in some subspace $V_k \subseteq \mathcal{N}_k(x)$ with $x' \in V_k$. When the local search used in this step is also VNS, the two-level VNS-scheme arises.

VNDS can be viewed as embedding the classical successive approximation scheme in the VNS framework.

8 Analyzing Performance

When a first VNS heuristic has been obtained and tested, the effort should not stop there. Indeed, it is often at this point that the most creative part of the development process takes place. It exploits systematically *Fact 2* of the Introduction, i.e., that global minima are local minima for all possible neighborhoods simultaneously. The contrapositive is that if a solution $x \in X$ is a local minimum (for the current set of neighborhoods) and not a global one there are one or several neighborhoods (or moves) to be found, which will bring it to this global optimum.

The study then focuses on instances for which an optimal solution is known (or, if none or very few are available, on instances with a presumably optimal solutions, i.e., the best one found by several heuristics) and compares it with the heuristic solution obtained. Visualization is helpful and make take the form of a *Distance-to-target diagram* (Hansen and Mladenović, 2003). Then, the heuristic solutions, the optimal one and their symmetric difference (e.g., for the Traveling Salesman Problem, TSP for short) are represented on screen. Moreover, an interactive feature allows a follow how the heuristic works step by step. The information thus gathered is much more detailed than what one would get just from objective values and computer times if, as is often the case, the heuristic is viewed as a black box. For instance, this clearly shows that 2-opt is not sufficient to get a good solution for the TSP, that moves involving 3 or 4 edges are needed and that those edges leading to an improvement may be far apart along the tour. For another application of VNS to the TSP see Burke *et al.* (1999).

Similarly, for location problems, one can focus on those facilities which are not at their optimal location and study why, in terms of distributions of nearby users.

Another point is to study how to get out of a large valley if there exists another promising one. *Valley* (or *mountain*) *profiles* are then useful (Hansen *et al.*, 2001). They are obtained by drawing many points x' at random within nested neighborhoods $\mathcal{N}_1(x), \mathcal{N}_2(x), \ldots$ (or, which is equivalent, at increasing distance of a local minimum x) then performing one VND descent and plotting probabilities to get back to x, to get to another local minimum x'' with a value $f(x'') \geq f(x)$ or to get to an improved local minimum x' with $f(x'') < f(x)$. Alternately one may also study the probabilities to go in the direction of x, i.e., $\rho(x, x'') \leq \rho(x, x')$ or towards another valley i.e., $\rho(x, x'') > \rho(x, x')$.

9   Promising areas of research

Research on Variable Neighborhood Search and its applications is currently very active. We review some of the promising areas in this section; these include a few which are barely explored yet.

A first set of areas concerns enhancements of the VNS basic scheme and ways to make various steps more efficient.

(a) **Initialization.** Both VND and VNS, as many other heuristics, require an initial solution. Two questions then arise: *How best to choose it?* and *Does it matter?* For instance, many initialization rules have been proposed for the k-Means heuristic for minimum sum-of-squares clustering, described above, 25 such rules are compared in Hansen *et al.* (2003). It appears that while sensitivity of k-Means to the initial solution is considerable (best results being obtained with Ward's hierarchical clustering method), VNS results depend very little on the chosen rule. The simplest one is thus best. It would be interesting to extend and generalize this result by conducting similar experiments for other problems.

(b) **Inventory of neighborhoods.** As mentioned above, a VNS study begins by gathering material on neighborhoods used in previous heuristics for the problem under study. A systematic study of moves (or neighborhoods) used for heuristics for whole classes of problems (e.g. location, network design, routing, ...) together with the data-structures most adequate for their implementation should be of basic interest for VNS as well as for other metaheuristics. Several researchers, e.g. Ahuja *et al.* (2000) are working in that direction.

(c) **Distribution of neighborhoods.** When applying a General VNS scheme, neighborhoods can be used in the local search phase, in the shaking phase or in both. A systematic study of their best distribution between phases could enhance performance and provide further insight in the solution process. In particular, the trade-off between increased work in the descent, which provides better local optima, and in shaking which leads to better valleys should be focussed upon.

(d) **Ancillary tests.** VNS schemes use randomization in their attempts to find better solutions. This also avoids possible cycling. However, many moves may not lead to any improvement. This suggests to add ancillary test (Hansen, 1974, 1975) the role of which is to decide if a move should be used or not, in its general or in a restricted form. Considering again minimum sum-of-squares clustering, one could try to select better the centroid to be removed from the current solution (a possible criterion being that its cluster contains a few entities only or is close to another centroid) as well as the position where it will be assigned (e.g., the location of an entity far from any other centroid and in a fairly dense region).

A second set of areas concerns changes to the basic scheme of VNS.

(e) Use of memory. VNS in its present form relies only on the best solutions currently known to center the search. Knowledge of previous good solutions is forgotten, but might be useful to indicate promising regions not much explored yet. Also characteristics common to many or most good solutions, such as variables taking the some value in all or most such solutions could be used to better focus the shaking phase. Use of memory has been much studied in Tabu search and other metaheuristics. The challenge for VNS would be to introduce memory while keeping simplicity.

An interesting way to use memory to enhance performance is *Reactive VNS*, explored by Braisy [**?**] for the Vehicle routing problem with time windows. If some constraints are hard to satisfy their violation may be penalized more frequently than for others in the solution process.

(f) Parallel VNS. Clearly, there are many natural ways to parallelize VNS schemes. A first one, within VND, is to perform local search in parallel. A second one, within VNS, is to assign the exploration of each neighborhood of the incumbent to a different processor. A third one, within VNDS, is to assign a different subproblem to each processor. Lopez *et al.* (2002) explore several options in designing a parallel VNS.

(g) Hybrids. Several researchers, e.g. Rodriguez *et al.* (1999), Festa *et al.* (2001), Ribeiro *et al.* (2001), Drezner (2003a, 2003b) have combined VNS with other metaheuristics for various problems. Again, this is not always easy to do without losing VNS' simplicity but may lead to excellent results particulary if the other metaheuristics are very different from VNS.

At a more general level one might wish to explore combinations of VNS with *constraint programming*, instead of its development within mathematical programming as in the applications described above. This could be done in two directions: on the one hand, techniques from constraint programming could be applied to enhance VND; on the other hand, VNS could be applied to constraint programming by minimizing a sum of artificial variables measuring infeasibility and possibly weighted by some estimate of the difficulty of satisfying the corresponding constraints.

A third set of areas concerns new aims for VNS, i.e., non-standard uses.

(h) Solutions with bounds on the error. VNS, as other metaheuristics, most often provides near-optimal solutions to combinatorial problems, without bounds on their error. So while such solutions may be optimal or very close to optimality, this fact cannot be recognized. One approach to obtain such bounds is to find with VNS a heuristic solution of the primal problem, deduce from it a solution to the dual (or its continuous relaxation) and then improve this dual solution by another application of VNS. Moreover, complementary slackness conditions can be used to simplify the dual. For problems with a small duality gap this may lead to near optimal solution

guaranteed to be very close to optimality. To illustrate, recent work of Hansen *et al.* (2003) on the Simple Plant Location Problem gave solutions to instances with up to 15,000 users and 15,000 possible facilities with an error bounded by 0.05%.

(i) **Using VNS within exact algorithms for mixed-integer programming.** Sophisticated algorithms for mixed-integer programming often contain various phases where heuristics are applied. This is illustrated e.g. by Desaulniers, Desrosiers and Solomon (2001) for the airline crew scheduling problem.

Extending the results described in the previous subsection, in the branch-and-bound framework led to solve exactly SPLP instances with up to 7,000 users (Hansen *et al.*, 2003).

A different approach, called *Local branching*, has been recently proposed by Fischetti and Lodi (2003) and Fischetti *et al.* (2003), both for exact and approximate resolution of large mixed-integer programs. At various branches in the branch-and-bound tree, cuts (which are not valid in general) are added; they express that among a given set of 0-1 variables, already at an integer value, only a few may change their value. They thus corresponds to neighborhoods defined by the Hamming distance. Then CPLEX is used to find the optimal solution within the neighborhood and in this way feasible solutions are more easily obtained. Improved solutions were obtained for a series of large mixed-integer programming instances from various sources.

(j) **Artificial intelligence: enhancing graph theory with VNS.** VNS, as other metaheuristics has been extensively used to solve a variety of optimization problems in graph theory. However, it may also be used to enhance graph theory *per se*, following an Artificial Intelligence approach. This is done by the AutoGraphiX (AGX) system developed by Caporossi and Hansen (2000, 2003). This system considers a graph invariant (i.e., a quantity defined for all graphs of the class under study and independent of vertex and edge labelings) or a formula involving several invariants (which is itself a graph invariant). Then AGX finds extremal or near-extremal graphs for that invariant parameterizing on a few variables, often the order n (or number of vertices) and the size m (of number of edges) of the graph. Analyzing automatically or interactively these graphs and the corresponding curves of invariant values leads to find new conjectures, refute, corroborate or strengthen existing ones and get hints about possible proof from the minimal list of moves needed to find the extremal graphs. To illustrate, the *energy E* of a graph is the sum of absolute values of the eigenvalues of its adjacency matrix. The following relations were obtained by Caporossi et al. (1999) with AGX: $E \geq 2\sqrt{m}$ and $E \geq \frac{4m}{n}$ and were easily proved. Over 70 new relations have now been obtained, in mathematics and in chemistry. Three ways to attain full automation based on the mathematics of principal component analysis, linear programming and recognition of extremal graphs together with formula manipulations are currently being studied.

10   Conclusions

The general schemes of Variable Neighborhood Search have been presented, discussed and illustrated by examples. References to many further successful applications are given in the Appendix. In order to evaluate the VNS research program, one needs a list of desirable properties of metaheuristics. The following eight ones are presented in Hansen and Mladenović (2003):

  (i) *Simplicity:* the metaheuristic should be based on a simple and clear principle, which should be largely applicable;

 (ii) *Precision:* steps of the metaheuristic should be formulated in precise mathematical terms, independent from the possible physical or biological analogy which was an initial source of inspiration;

(iii) *Coherence:* all steps of heuristics for particular problems should follow naturally from the metaheuristic's principle;

 (iv) *Efficiency:* heuristics for particular problems should provide optimal or near-optimal solutions for all or at least most realistic instances. Preferably, they should find optimal solutions for most problems of benchmarks for which such solutions are known, when available;

  (v) *Effectiveness:* heuristics for particular problems should take moderate computing time to provide optimal or near-optimal solutions;

 (vi) *Robustness:* performance of heuristics should be consistent over a variety of instances, i.e., not just fine-tuned to some training set and less good elsewhere;

(vii) *User-friendliness:* heuristics should be clearly expressed, easy to understand and, most important, easy to use. This implies they should have as few parameters as possible and ideally none;

(viii) *Innovation:* preferably, the metaheuristic's principle and / or the efficiency and effectiveness of the heuristics derived from it should lead to new types of applications.

As argued there, as well as in the more recent among the surveys listed in the Appendix, VNS possesses, to a large extend, all of those properties. This has led to heuristics among the very best ones for several problems, but more importantly to insight into the solution process and some innovative applications.

Appendix. Sources of Additional Information about VNS

Some of Web addresses with sources of information about Variable Neighborhood Search include

- **http://www.mi.sanu.ac.yu/VNS/VNS.HTM**
    This is a working web presentation of VNS, developed by Tatjana Davidović, PhD student at University of Belgrade;

- **VNSHeuristic.ull.es**
    Another web page for the VNS designed by Professor's Moreno research group from University of La Laguna.

- **http://www.gerad.ca/en/publications/cahiers.php**
    If one choose the option "search for papers" and in the "Abstract" box type "Variable Neighborhood Search", 23 papers for downloading will appear at the screen;

- **http://smg.ulb.ac.be**
    There are several papers on VNS in "Preprints" by Hansen, Labbé, Mélot, Mladenović, etc.

Survey papers are: P. Hansen and N. Mladenović (1999, 2001a, 2001c, 2002a, 2002b, 2003), Hansen, Mladenović and Moreno-Perez (2003), and Kochetov, Hansen and Mladenović (2003).

References

R.K. Ahuja, J.B. Orlin and D. Sharma (2000). Very large-scale neighborhood search, *International Transactions in Operations Research*, 7: 301–317.

D.J. Aloise, D. Aloise, C.T.M. Rocha, J.C. Ribeiro Filho, L.S.S. Moura and C.C. Ribeiro (2003). Scheduling workover rigs for onshore oil production, Research Report Dpt. of Comp. Sci., Catholic University of Rio de Janeiro (submitted).

C. Avanthay, A. Hertz and N. Zufferey (2003). Variable neighborhood search for graph coloring *European J. Opernl. Res.* 151: 379–388

E.B. Baum (1986). Toward practical 'neural' computation for combinatorial optimization problems. In J. Denker (ed.), *Neural networks for computing*, American Institute of Physics.

O. Braysy (2001). *Local search and variable neighborhood search algorithms for vehicle routing with time windows*, Acta Wasaensia 87, Universitas Wasaenis, Vaasa.

E.K. Burke, P. Cowling and R. Keuthen (1999). Effective local and guided Variable neighborhood search methods for the asymmetric traveling salesman problem. In the Proceedings of the Evo Workshops, Springer, Lecture Notes in Computer Science, pp. 203–212.

G. Caporossi, D. Cvetković, I. Gutman and P. Hansen (1999). Variable neighborhood search for extremal graphs. 2. Finding graphs with extremal energy, *J. Chem. Inf. Comput. Sci.*, 39: 984–996.

G. Caporossi and P. Hansen (2000). Variable neighborhood search for extremal graphs 1. The AutoGraphiX system. *Discrete Mathematics*, 212: 29–44.

G. Caporossi and P. Hansen (2003). Variable neighborhood search for extremal graphs 5. Three ways to automate conjecture finding, *Discrete Mathematics* (to appear).

G. Cornuejols, M. Fisher and G. Nemhauser (1990). The uncapacitated facility location problem. In P. Mirchandani and R. Francis (eds), *Discrete Location Theory*, John Wiley & Sons.

G. Desaulniers, J. Desrosiers and M.M. Solomon (2001). Accelerating strategies in column generation methods for vehicle routing and crew scheduling problems. In: *Essays and Surveys in Metaheuristics*, pp. 309–324, Kluwer Academic Publishers, Boston/Dordrecht/London.

Z. Drezner (2003a). Heuristic algorithms for the solution of the quadratic assignment problem, *J. Appl. Math. and Decision Sci.*, (to appear).

Z. Drezner (2003b). A new genetic algorithm for the quadratic assignment problem, *INFORMS J. on Computing* 15 (to appear).

P. Festa, P. Pardalos, M. Resende and C. Ribeiro (2001). GRASP and VNS for Max-cut, Proceedings of MIC'2001, pp. 371–376.

M. Fischetti and A. Lodi (2003). Local branching. Mathematical Programming B, Published on line, March 28.

M. Fischetti, C. Polo and M. Scantamburlo (2003). A local branching heuristic for mixed-integer programs with 2-level variables, Research report, University of Padova.

M.R. Garey and D.S. Johnson (1978). *Computers and Intractability: A Guide to the Theory of NP-Completeness*, Freeman, New-York.

F. Glover and M. Laguna (1997). *Tabu search*. Kluwer, Boston.

F. Glover and G. Kochenberger (Eds.) (2003). *Handbook of Metaheuristics*, Kluwer.

P. Hansen. *Programmes mathématiques en variables 0-1* (1974). Thèse d'Agrégation de l'Enseignement Supérieur, Université Libre de Bruxelles.

P. Hansen (1975). Les procédures d'optimization et d'exploration par séparation et évaluation, *Combinatorial Programming*, (in B. Roy ed.) pp. 19–65, Dordrecht: Reidel.

P. Hansen, J. Brimberg, D. Urošević and N. Mladenović (2003). Primal - dual variable neighborhood search for exact solution of the simple plant location problem (in preparation).

P. Hansen, B. Jaumard, N. Mladenović and A. Parreira (2000). Variable neighborhood search for weighted maximum satisfiability problem, *Les Cahiers du GERAD* G-2000-62, Montréal, Canada.

P. Hansen and N. Mladenović (1997). Variable neighborhood search for the p-Median, *Location Sci.*, 5: 207–226.

P. Hansen and N. Mladenović (1999). An introduction to variable neighborhood search. In: S. Voss *et al.* (eds), *Metaheuristics, Advances and Trends in Local Search Paradigms for Optimization*, pp. 433–458, Kluwer, Dordrecht.

P. Hansen and N. Mladenović (2001a). Variable neighborhood search: Principles and applications. *European J. of Oper. Res.*, 130: 449–467.

P. Hansen and N. Mladenović (2001b). J-means: A new local search heuristic for minimum sum-of-squares clustering, *Pattern Recognition*, 34: 405–413.

P. Hansen and N. Mladenović (2001c). Developments of variable neighborhood search. In: C. Ribeiro and P. Hansen (eds.), *Essays and Surveys in Metaheuristics*, pp. 415–440, Kluwer Academic Publishers, Boston/Dordrecht/London.

P. Hansen and N. Mladenović (2002a). Variable neighborhood search. In: P. Pardalos and M. Resende (eds), *Handbook of Applied Optimization*, pp. 221–234, Oxford University Press, New York.

P. Hansen and N. Mladenović (2002b). Recherche à voisinage variable. In: J. Teghem and M. Pirlot (eds), *Optimisation approchée en recherche opérationnelle*, pp. 81–100, Lavoisier, Hermès Science Publications, Paris.

P. Hansen and N. Mladenović (2003). Variable neighborhood search. In: F. Glover and G. Kochenberger (eds), *Handbook of Metaheuristics*, pp. 145–184, Kluwer Academic Publisher.

P. Hansen, N. Mladenović and J. A. Moreno Pérez (2003). Búsqueda de Entorno Variable (in Spanish), *Inteligencia Artificial* (to appear).

P. Hansen, N. Mladenović and D. Perez-Brito (2001). Variable neighborhood decomposition search. *J. of Heuristics*, 7(4): 335–350.

P. Hansen, E. Ngai, B. Cheung and N. Mladenović (2003). Survey and comparison of initialization methods for k-means clustering (in preparation).

S.C.D Kirkpatrick, Jr. Gellatt, and P. Vecchi (1983). Optimization by Simulated Annealing. *Science*, 220: 671–680.

Y. Kochetov, N. Mladenović and P. Hansen (2003). Lokalnii poisk s chereduyshimisy okrestnostyami (in Russian), *Diskretnoi analiza* (to appear)

M. Labbé, D. Peeters and JF Thisse (1995). Location on networks. In: Ball, M. *et al.* (eds), *Network Routing*, pp. 551–624, North-Holland, Amsterdam.

F.G. Lopez, B.M.Batista, J.A.Moreno Pérez and J.M.Moreno Vega (2002). The parallel variable neighborhood search for the p-median problem. *J. of Heuristics* 8: 375–388.

N. Mladenović and P. Hansen (1997). Variable neighborhood search. *Computers Oper. Res.* 24: 1097–1100.

C. Papadimitriou (1994). *Computational Complexity*, Addison Wesley.

C.R. Reeves, (Eds.) (1992). *Modern Heuristic Techniques for Combinatorial Problems*. Blackwell Scientific Press, Oxford, UK.

M.G.C. Resende nad R. Werneck (2003). On the implementation of a swap-based local search procedure for the p-median problem. In: Proceedings of the *Fifth Workshop on Algorithm Engineering and Experiments* (ALENEX'03), R.E. Ladner (ed.), SIAM, Philadelphia, pp. 119-127.

C. Ribeiro, E. Uchoa and R. Werneck (2001). A hybrid GRASP with perturbations for the Steiner problem in graphs. Technical report, Computer Science Department, Catholic University of Rio de Janeiro.

I. Rodriguez, M. Moreno-Vega and J. Moreno-Perez (1999). Heuristics for routing-median problems. SMG Report, Université Libre de Bruxelles, Belgium.

M.B. Teitz and P. Bart (1968). Heuristic methods for estimating the generalized vertex median of a weighted graph. *Operations Research*, 16(5): 955–961.

R. Whittaker (1983). A fast algorithm for the greedy interchange for large-scale clustering and median location problems. *INFOR* 21: 95–108.